CLMM Aksel Kretsinger-Walters, adk2164
Hopfield Networks is All You Need adk2164@columbia.edu
COMS 6998: Spring 2026 **Due:** 03/25/2026 11:59 PM ET

1 Hopfield Networks is All You Need (Ramsauer et al., 2020)

1.1 Motivation

Classical Hopfield networks are associative memory models that store patterns and retrieve them via energy minimization. However, traditional formulations suffer from:

- Limited storage capacity
- Shallow energy landscapes
- Difficulty scaling to modern machine learning settings

This paper revisits Hopfield networks and shows that a modernized formulation can scale to high-dimensional settings and is closely related to attention mechanisms.

1.2 Key Idea

The paper introduces **modern Hopfield networks** with continuous states and an updated energy function. The central insight is:

> Modern Hopfield updates are mathematically equivalent to attention mechanisms.

This connects associative memory with transformer architectures.

1.3 Model Formulation

Given stored patterns $\{x_i\}_{i=1}^N$, the network defines an energy function:

$$E(\xi) = -\frac{1}{\beta} \log \left(\sum_{i=1}^{N} \exp(\beta x_i^T \xi) \right)$$

The update rule becomes:

$$\xi \leftarrow \sum_{i=1}^{N} \text{softmax}(\beta x_i^T \xi) \; x_i$$

This is equivalent to a weighted sum of stored patterns using a softmax over similarities.

1.4 Connection to Attention

This update rule is identical to attention:

$$\text{Attention}(q, K, V) = \text{softmax}(qK^T)V$$

Mapping:

- Query $q \leftrightarrow \xi$
- Keys $K \leftrightarrow x_i$
- Values $V \leftrightarrow x_i$

Thus, attention can be interpreted as a retrieval step in a Hopfield network.

1.5 Memory Properties

Modern Hopfield networks have dramatically improved capacity:

- Exponential storage capacity in dimension
- Ability to store continuous patterns
- Fast convergence (often one update step)

They can retrieve:

- Exact stored patterns
- Averaged patterns
- Global prototypes

1.6 Energy Landscape

The energy function induces different types of fixed points:

- Individual memories (sharp minima)
- Metastable states (clusters)
- Global averages (broad minima)

This explains why attention can interpolate between patterns.

1.7 Applications

The paper demonstrates that modern Hopfield layers can be used in:

- Multiple instance learning
- Set-based learning tasks
- Deep neural networks as a drop-in module

They provide a principled interpretation of attention as memory retrieval.

1.8 Key Contributions

- Reformulation of Hopfield networks with continuous states
- Proof of equivalence between Hopfield updates and attention
- Exponential increase in storage capacity
- Integration of associative memory into deep learning

1.9 Takeaway

The paper shows that:

$$\text{Attention} \approx \text{Modern Hopfield Retrieval}$$

This unifies classical associative memory with modern transformer architectures and provides a theoretical foundation for attention mechanisms.