CLMM
Continual Learning Overview: Summary
COMS 6998: Spring 2026

Aksel Kretsinger-Walters, adk2164
adk2164@columbia.edu
Due: 01/28/2026 11:59 PM ET

Nested Learning

Motivation. Modern machine learning systems combine neural architectures and optimization algorithms in complex ways. Despite advances in deep learning, there remains limited understanding of how learning algorithms, architectures, and memory mechanisms interact.

The Nested Learning (NL) paradigm proposes a unifying framework that models learning systems as hierarchies of interconnected optimization processes.

Core Idea. Nested Learning represents a machine learning model as a collection of **nested optimization problems**, each operating at a different level of abstraction. [**?**]

In this view:

- Neural architectures operate on **data context** (tokens, images, etc.).
- Optimizers operate on **gradient context**.

Both processes can be interpreted as learning systems that compress their respective contexts.

This perspective suggests that optimization algorithms and neural architectures are fundamentally the same type of object operating at different levels.

Associative Memory Interpretation. Nested Learning interprets many machine learning components as forms of associative memory.

For example:

- Neural networks compress training data into parameters.
- Optimizers compress gradient information into momentum or adaptive statistics.

Under this interpretation, common optimizers such as Adam or SGD with momentum are themselves associative memory systems that learn to store gradient information.

Multi-Level Learning Systems. The NL framework introduces multiple interacting learning levels:

(a) Data-level learning (standard neural network training)

(b) Optimizer-level learning (learning how to update parameters)

(c) Meta-learning levels that adapt the learning algorithm itself

Each level maintains its own internal gradient flow and update rule.

Self-Modifying Models. Building on this framework, the paper proposes architectures that can modify their own learning rules. These models learn update rules that control how their internal memory and parameters evolve over time.

This enables forms of continual learning and adaptive reasoning.

Continuum Memory System. The paper also introduces a *continuum memory system* in which memory exists along a spectrum of update frequencies rather than being divided into discrete short-term and long-term components.

Different memory modules update at different temporal scales, enabling both rapid adaptation and persistent knowledge storage.

Hope Architecture. As a proof of concept, the authors introduce a continual learning model called *Hope*, which integrates:

- nested optimization mechanisms
- adaptive memory modules
- self-modifying learning rules

Experiments demonstrate promising results on language modeling, continual learning benchmarks, and few-shot generalization tasks.

Limitations. The Nested Learning framework is primarily conceptual and theoretical. Many proposed mechanisms remain difficult to implement efficiently in large-scale models.

Additionally, the framework introduces substantial architectural complexity and raises challenges for stable training.

Takeaway. Nested Learning proposes a new perspective on deep learning in which architectures, optimizers, and memory systems are viewed as nested learning processes. This unified framework may provide insights for designing more adaptive and continual learning systems.